

PhD supervision meeting — Tue 23 Jun 2026, 1:00–4:00 pm (online)

Single source for the meeting. Replaces the old brief / primer / cheatsheet split. Two adjacent projects, two opposite probes, one consolidated playbook. Numbers flagged (*approx*) come from secondary summaries — say “roughly” if you quote them.

Mindset before you dial in — read this last (it overrides the tactics below)

Everything below is your *private map*, not a script to perform. A first supervision meeting is a mutual audition for a 4-year working relationship, not an interview you administer. Walk in as an **enthusiast, not a prosecutor**.

- **Ask them first.** Open with “*What’s the question in this project you find most exciting — where do you see it going?*” Let them talk. You’ll hear the predict-vs-explain answer naturally, no probe required.
- **Lead with ONE idea you’re genuinely excited about**, not a decision tree. Specific, feasible, adjacent to their work. Enthusiasm + a concrete first study is what gets a candidate funded.
- **Show competence, don’t assert it.** You have **no NLP yet** — the supervisor’s first worry. Get ahead of it: run a few VN earnings headlines through GPT-4 vs a **Vietnamese baseline (PhoBERT-class — VADER is English-only)** this weekend and mention what you saw. A tiny real demo beats any “I ramp fast.”
- **Hold the structural asks** (DR203 route, grafting Anh on). Raising enrolment-code mechanics early signals you’re gaming classification rather than burning to do the research. Earn them; raise later, possibly over email.
- **Keep these words OUT of the room:** “gatekeeper,” “fallback,” “house-gravity,” “probe.” Sharp private strategy, corrosive public posture.
- **Evaluate the humans, not just the projects.** Can you take hard feedback from Buerthey for years? Will Anh actually be engaged? You’re choosing people, not a topic.
- **A redirect you can’t get is information, not defeat.** If an accounting lead genuinely can’t supervise an after-cost backtest, that’s a real answer — receive it graciously, don’t push.

Demote this whole document from a script you execute to a map you’ve memorised. If you’re listening for your branch triggers, you’re not actually listening.

0. The one thing to learn

Is the thesis judged on **predicting returns out-of-sample, after costs** — or on **explaining** a relationship? That single answer routes the whole meeting. You ask the *neutral* version of the question; the **answer** is the diagnostic — never ask the diagnostic out loud.

1. The room & the situation

In the room (all cc'd on Buerthey's thread): - **Samuel Buerthey** — DR200 (PhD Accountancy) lead. Accounting / ESG-disclosure (h-12, ~100% CSR/disclosure, **0 asset-pricing**). *Topic gatekeeper on the PEAD arm.* - **Anh Tuan Nguyen** — Lecturer in Finance, RMIT VN (PhD Finance, UT Arlington; REIT/property mispricing, fund flows, market efficiency). **Your method ally on EITHER project.** Junior output (h-3, ~26 cites) — pin alpha framing to him, but he can't anchor a thesis alone. *Co-author of the ChatGPT-Bitcoin paper.* - **Hà Xuân Sơn** — Lecturer, Blockchain Enabled Business (~1.7k cites, **CS/security, not finance**). **Also leads the LLM project #192.** Brings the AI/LLM engine, not asset-pricing judgment.

This is a DUAL-PURPOSE meeting. It opens as the PEAD (DR200) conversation. But the team's own published work is in-sample/explanatory, so the ~80%-likely answer to the probe is "explanatory" — at which point **this becomes the LLM #192 meeting.** For the LLM pivot, **Son (a #192 co-lead) and Anh (the finance co-sup you'd want grafted on) are both already in the room** — though **Binh, #192's senior co-lead, is not**, so a firm #192 commitment may need his later sign-off. So **both arms are prepped to the same depth.**

Goal: find the right *home* for ONE coherent research interest across two adjacent projects. - **Target:** LLM #192 enrolled under **DR203 (Economics, Finance & Marketing)** + a finance co-sup. - **Fallback:** PEAD #247 / DR200 — if its framing can be made genuinely predictive. - **Leave with PEAD intact either way.** Don't burn the fallback.

Why this configuration — neither project passes *both* prongs of "predictive verbs AND finance discipline/code/supervisor": - **LLM #192** — brief is alpha-native ("predict asset prices", benchmarks vs ARIMA/factor models, "LLM-based trading systems") ☒, and **DR203 (Finance)** is **already one of its listed enrolment codes** ☒ ... but FoR is 0% finance, discipline is Business (GSBL), and there's **no named asset-pricing supervisor** ✕. - **PEAD #247** — has a finance co-sup (Anh) + a 30% finance FoR share ☒ ... but framing is accounting-EXPLANATORY ("explain over/under-reaction"), and it's **locked to DR200 (Accountancy)** — **no finance enrolment route** ✕.

Starting from the alpha-native brief and *adding* a supervisor beats starting from an explanatory brief and *fighting* the framing past an accounting lead.

1.5 The two projects, side by side

Dimension	LLM #192 (<i>target</i>)	PEAD #247 (<i>fallback</i>)
Full title	Large Language Models for Business Innovation and Decision-Making	Investor Sentiment, Earnings Surprises & Market Anomalies Leveraging Generative AI in Capital Market Research
Enrolment code	DR201 / DR205 / DR203 — DR203 (Finance) already a listed route	DR200 only — PhD (Accountancy); no finance-code route
Discipline	Business (GSBL) — generic, flexible	Accounting (AISSC)
FoR weight-ing	40% Org Behaviour · 35% Data Mining · 25% Innovation Mgmt — 0% finance	40% Financial Accounting · 30% Finance · 30% AI
Lead team	Hà Xuân Sơn (CS) + AP Binh Nguyen (senior; econometrics/crypto, h-11) — no asset-pricing lead	Buertey (accounting) + Anh Nguyen (finance) + Ha Son (AI)
Funding closes	· “Funded” (full-time only — likely <i>not</i> part-time; see §11) · 30 Jun 2027	“Funded” (same caveat) · 31 Dec 2027
Framing verbs	“predict asset prices”, “informational efficiency”, “LLM-based trading systems ”	“explain over/under-reaction”, “amplifies or mitigates PEAD”, “price discovery”
ESG check	ECP “Design & Creative Practice”; no SDG tag	ECP “Information in Society”; no SDG tag → not an ESG project

The decisive asymmetry: #192 *already offers DR203 (Finance) enrolment* but has no finance supervisor and 0% finance FoR; #247 has a finance supervisor and a 30% finance FoR but is locked to DR200 (Accountancy). So the LLM path’s only genuinely missing piece is a **finance co-supervisor (Anh)** — the finance *enrolment* is already on the table. That’s a lighter lift than importing predictive framing past an accounting lead.

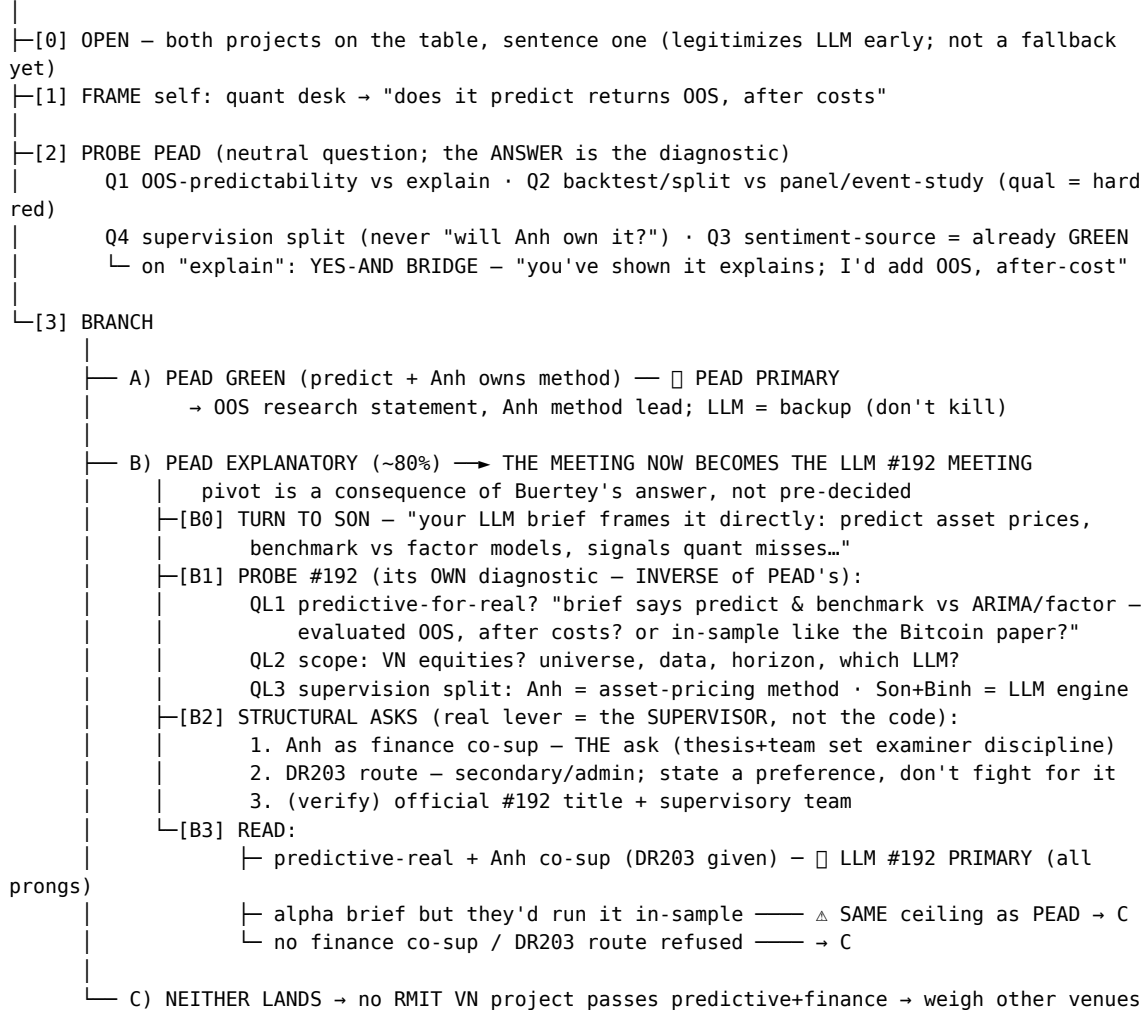
Phrases to mirror back (signals you read the briefs): - #192: “predict asset prices, characterise investor behaviour, and assess the **informational efficiency** of financial markets”; signals “traditional quantitative models miss”; the reflexivity hook — “*does AI adoption change the very market dynamics it seeks to predict?*”; relevance to “investors, regulators, and policymakers... market structure and financial stability.” - #247: “whether investor sentiment **amplifies or mitigates** anomalies such as PEAD”; sentiment from “social media and online news... complementing Google Trends or media coverage”; “how sentiment, attention, and accounting information **jointly influence price discovery.**”

2. Decision tree

DUAL-PURPOSE: opens as PEAD/DR200; on the ~80%-likely "explanatory" answer it BECOMES the LLM #192 meeting.

In the room: Buerthey + Anh (PEAD) and Son (a #192 co-lead) – prep both arms equally.

Note: Binh, #192's senior co-lead, is NOT present, so a firm #192 commitment may need his sign-off.



RAILS (every branch): never voice "fallback"/"flawed"/the FoR-code critique · address asset-pricing to Anh ·

leave PEAD intact · cite the ChatGPT-Bitcoin paper TO ANH & SON (their paper, NOT Buerthey's)

3. Logistics

- They pick the time inside the **1:00–4:00 pm** window. CV already sent (docs/cv/quy-sang-mai-cv.pdf).
 - Fit line stays honest: quant / markets side — **not** claiming NLP / generative-AI experience yet.
-

4. SAY (in order)

[pre-0] 20-second intro (if they open with “tell us about yourself”) > “I’m a quant trader at HSC in Vietnam — CFA, MSc Applied Mathematics. My capstone built boosting / expert-weighting online-learning signals on VN equities, and I’d like to take that into LLM-driven sentiment signals for a PhD.”

[0] Open — shared ground, legitimise both projects > “Both your DR200 sentiment project and Son’s LLM project sit on the same ground I care about — sentiment, LLMs, earnings reactions, market anomalies. I wanted to use this to figure out which is the right home for the question I’d actually want to answer.”

[1] Frame myself (honest, QR) > “My background’s a quant trading desk, so I instinctively frame a sentiment question as ‘does this predict returns out-of-sample, after costs.’ I want to check that fits how you’d build the thesis.”

[2] THE decisive question — neutral either/or, ask early > “Would the thesis be judged on whether the sentiment signal **predicts returns out-of-sample, after costs**, or on **explaining** investors’ over/under-reaction to earnings? I want to build toward the right one.”

[2b] Supervision split — gets method-ownership WITHOUT asking for it > “Since it spans accounting, finance and AI across the three of you — how would supervision divide up? Who would I mainly work with on the empirical / asset-pricing side, versus the sentiment-extraction and earnings-surprise pieces?” > → **Let Buertey describe it first**. Anh handed the empirical lead = confirmed.

5. PEAD arm — read the answers

#	Ask	☑ GREEN (predictive)	× RED (explanatory/accounting)
Q1	Judged on OOS return predictability , or explaining a relationship?	“predict / survive costs / tradable”	“establish / explain the channel”
Q2	Predictive backtest w/ train-test split, or panel/event-study?	“walk-forward / hold-out / net-of-cost”	“panel / fixed effects” — qual = hard red flag
Q4	(<i>diagnostic — never ask verbally</i>) Does Anh own the empirical method?	“you’d work with Anh on the empirical side”	“I’d guide all the empirical framing myself”

(Q3 — sentiment source — is already GREEN: the brief specifies broad social media + news, **not** ESG/sustainability disclosure. Don’t relitigate it.)

The instant you hear “explain” → YES-AND BRIDGE (don’t act surprised — you already know this is likely): > “You’ve shown sentiment *explains* returns; I’d *add* the out-of-sample,

after-cost predictive test on VN equities — same data, same LLM sentiment layer, one extra evaluation lens.”

Their Bitcoin paper’s own illustrative Sharpe element is the hook that makes this land. The gap (no OOS / no after-cost) **is** your thesis.

6. LLM #192 arm — read the answers

The inverse problem: #192’s *brief* is already predictive, but the team’s *house style* is explanatory. So you confirm they’d actually build it predictively, then import the missing finance prong.

#	Ask	☒ GREEN	⚠ RED (same ceiling as PEAD)
QL1	Evaluated OOS, after costs, vs ARIMA/factor benchmarks — or in-sample like the Bitcoin paper?	“walk-forward / hold-out / net-of-cost / beat the factor model”	“contemporaneous / in-sample / explain the relationship”
QL2	Scope — VN equities? universe, data source, horizon, which LLM?	concrete, tradable framing	“we’d see / it’s open” with no markets realism
QL3	Supervision split — who owns the asset-pricing method?	“Anh on the empirical/finance side”	“Son and I would shape it” (no finance lead)

Structural asks (the prong #192 lacks) — float, don’t demand. The real lever is the SUPERVISOR, not the code. 1. **Finance co-supervisor (Anh) — the genuinely missing piece, and the real ask.** #192’s roster (Son CS + Binh crypto/econ) has **no asset-pricing voice**. *Correction to earlier framing:* what makes a thesis “examined as finance” is the **thesis content + the team’s examiner nominations**, NOT the program code — so getting a finance supervisor onto the team is the lever that actually matters. (Anh likely can’t be the *registered primary* supervisor — too junior; aim for **Binh or another senior as principal, Anh as associate** co-sup.) 2. **DR203 enrolment — secondary/admin, not a fight.** DR203 is already a listed code, but choosing among a project’s codes is an *admissions/School* decision, not yours, and **the code doesn’t bind examiners**. Express a preference for the finance route; don’t treat it as the win. Confirm availability for the next intake — better over email than pressed in the room. 3. **Verify the basics** — “#192” is an internal tracking number; not in RMIT’s public HDR database. Ask the official **title and full supervisory team** to confirm.

Lead with the HYPOTHESIS, not the method (the identification angle — your strongest, most original framing): > “Vietnam has **no equity short-selling** and **±7% daily price limits**, so limits-to-arbitrage predict that post-earnings underreaction should be *larger and more persistent* than in the US — but also partly *uncapturable* on the long-only side. I want to measure how much of that residual drift an LLM sentiment signal captures, net of cost, for a long-only / VN30-

futures investor.” > → This turns VN’s market frictions from *caveats* into the **identification of the contribution** — and only someone with desk knowledge can pose it. (See §7.3 emerging-market PEAD.)

Scope answer — have this ready for “what would you actually build?”: > “VN30 / liquid-HOSE universe; LLM sentiment from Vietnamese financial news + earnings announcements — a **Vietnamese-language model (PhoBERT-class), not an English tool**; predict post-announcement drift returns. Out-of-sample **conditional on a verifiable LLM training cutoff** (test window strictly post-cutoff, robustness across model vintages), net of ~0.2–0.5% round-trip **plus a market-impact model (not a flat cost)** including the 0.1% sell tax. **Long-only or via VN30 futures — no real stock short-selling yet.** Control for ±7% HOSE price-limit censoring and the FTSE upgrade (21 Sep 2026) flow confound. Point-in-time FiinPro data. Benchmarks **by role**: a no-predictability null, rival ML predictors (LSTM/GBRT on the same features, with vs without LLM sentiment), and a VN-appropriate factor model for risk-adjustment — *not ARIMA*, which isn’t a meaningful null for a cross-sectional drift signal.” (See §7.4 and §7.7.)

7. Substance bank (shared across both arms)

7.1 The team’s OWN paper — your single best read

“Decoding market sentiment: the power of ChatGPT in explaining Bitcoin returns from X data” — *China Finance Review International*, 2025. DOI 10.1108/CFRI-05-2024-0278. **Authors: Binh Thanh Nguyen, Tuan Thanh Chu, Son Xuan Ha (= Hà Xuân Sơn), Anh Tuan Nguyen** — all RMIT VN. **△ Buerthey is NOT a co-author.** This binds Anh + Son + Binh to your exact topic (Son + Binh lead #192; Anh is the finance graft). **Cite it to Anh and Son**, never as “Buerthey’s group’s work.” - **Method:** zero-shot prompting — clean tweets → ChatGPT-4o/3.5 classifies Bullish/Bearish/Neutral → Bullishness + Agreement indices. Benchmarked vs BERT and VADER. No fine-tuning. - **Data:** crypto-influencer tweets + BTC daily close (Bitstamp), 3 Jan 2018 – 13 Jun 2023; controls = Fear&Greed, Google Trends, Wikipedia views. - **Explain vs predict: EXPLANATORY / in-sample** — contemporaneous (same-day) sentiment-on-returns regressions, heavy controls, BERT/VADER benchmarks. **No train/test split, no OOS window, no after-cost backtest.** The title’s “explaining” matches the evaluation exactly. - **Use it as:** proof the topic is live AND proof of the open door — “*you’ve shown sentiment explains returns; I’d make the Sharpe element rigorous: VN equities, genuine OOS, after-cost.*”

7.2 The predictive template — Lopez-Lira & Tang (2023)

“Can ChatGPT Forecast Stock Price Movements?” arXiv:2304.07619 / SSRN 4412788 (forthcoming JFE). The alpha-native twin of the team’s paper — cite as the model for your predictive chapter. - News headlines → ChatGPT (good/bad/unknown) → score → predict **next-day** returns; tested only on **post-training-cutoff** data = genuinely OOS. - Long-short predicts the drift; alpha concentrates in **small-caps + negative news**; rises with model size (GPT-4 > GPT-3.5; BERT/GPT-2 can’t); **decays as LLM adoption spreads**. - **THE caveat (have it ready):** cost/turnover figures (*approx, secondary — verify before quoting*) suggest it struggles at realistic costs; alpha sits exactly where costs bite hardest. So the honest framing is *capacity-aware, after-cost* evaluation. -

Glasserman & Lin (2023): look-ahead bias is *not* the dominant driver; LLM company knowledge can even hurt (“distraction effect”). - **The literature is unsettled** — the 2025 “*New Quant*” survey (arXiv:2510.05533, synthesising 50+ studies) reads LLM value as **conditional/complementary** (events, regime shifts); there’s **no clear evidence LLMs durably beat ARIMA/LSTM/factor baselines OOS net of costs** (your characterisation of the state of play, not a headline claim of the paper). Foreground this honesty; it disarms a “but is it tradable?” push.

7.3 PEAD canon + “is it still tradable?” (table stakes with an accounting lead)

- **Ball & Brown (1968)**: prices move with the earnings surprise — founding capital-markets-accounting paper.
- **Bernard & Thomas (1989, JAR / 1990, JAE)**: the drift — abnormal returns keep drifting ~60 trading days post-announcement; investors **underreact**. The anchor PEAD cite.
- **Debate**: risk premium (efficient) vs behavioral underreaction (mispricing) — unsettled; frame as open, don’t pick a side.
- **Emerging-market PEAD (don’t cite only US work)**: the drift is documented across Asian/emerging markets — e.g. Griffin, Kelly & Nardari (2010, RFS) on EM efficiency, plus China/Korea/Taiwan PEAD studies. An examiner *will* ask “why should US PEAD magnitudes hold in a frontier market with price limits, retail-dominated flow, and no shorting?” VN’s frictions could make underreaction *larger* (more mispricing) yet partly *uncapturable* — which is exactly the limits-to-arbitrage hypothesis to test (see §6 identification angle). Have the EM evidence ready.
- **Still tradable?** McLean & Pontiff (2016, JF): anomalies ~58% lower post-publication (and ~26% lower out-of-sample even pre-publication), worst for low-liquidity/high-idio-risk predictors = PEAD’s exact profile. Chordia et al. (2009, FAJ): long-short PEAD ~0.04%/mo in the most-liquid stocks vs ~2.43%/mo in the most-illiquid; **transaction costs eat ~70–100%** of the profit.
- **Ready answer**: “Naive drift capture is largely arbitrated away — McLean–Pontiff show ~half gone post-publication, Chordia shows the rest concentrates in illiquid names where costs eat 70–100%. So the honest claim isn’t ‘PEAD prints money’ — it’s conditional residual alpha: net-of-cost, liquidity-screened, combined with orthogonal signals like LLM sentiment. A measurement-and-signal-combination thesis, not a free-lunch one.”

7.4 Benchmark fluency (for the LLM arm) — benchmarks by ROLE, not one ladder

An examiner will pounce if you list ARIMA/FF/LSTM as a single ladder; they test three different things: - **The null** = no cross-sectional predictability (a naive no-signal baseline). **ARIMA is NOT this** — it’s a univariate model of a *return series*, not a benchmark for a cross-sectional, event-conditional drift. Including it is benchmark theatre; drop it or be ready to justify “ARIMA on *what* series, and why is that the null?” - **Rival predictors** = LSTM/GRU/GBRT/elastic-net on the **same feature set** (with vs without LLM sentiment) — the apples-to-apples ML competitor (Gu-Kelly-Xiu 2020, RFS). *Caveat*: many LSTM “wins” are on price *levels* (trivially autocorrelated), not returns. - **Risk adjustment** = Fama-French / characteristic controls applied to whatever you trade (significant intercept = unexplained alpha). Note: VN factors built on ~30 liquid names are **fragile** — flag it, don’t oversell.

7.5 Vietnam-equities realities (your edge — shows markets realism)

- **Exchanges:** HOSE (liquid large caps; VN-Index, VN30 = the investable universe), HNX (smaller), UPCoM (largely paper-only for backtests).
- **Price limits censor returns / block limit-up fills:** HOSE $\pm 7\%$, HNX $\pm 10\%$, UPCoM $\pm 15\%$.
- **No real stock short-selling yet** → long-short untradeable on the short leg; use **long-only** or **VN30 futures**. (T+0 / SBL roadmapped 2026–28.)
- **Foreign Ownership Limit “two-price” problem:** at full FOL, foreigners transact at an off-screen premium on top blue chips (banks, VCB, FPT) — quoted price $\neq$ achievable foreign price. State your investor assumption. (Circular 68, Nov 2024, removed foreign pre-funding.)
- **Regime breaks to control for:** KRX system live on HOSE (5 May 2025); **FTSE Frontier** → **Secondary-Emerging upgrade effective 21 Sep 2026** (~28 stocks get passive inflows — a demand-shock confound). MSCI still Frontier (*re-verify*).
- **Costs:** round-trip ~0.2–0.5% incl. mandatory 0.1% sell tax (charged even at a loss). T+2 settlement.
- **Data:** FiinGroup/FiinPro = academic standard (best point-in-time integrity). vnstock is convenient but survivorship/PIT-prone. Traps: survivorship (cross-exchange ticker migration), restatement/PIT bias, short history, stock-dividend back-adjustment, stale prices in thin names.

7.6 The reflexivity angle — a distinctive #192 chapter

#192’s brief explicitly asks: “*does AI adoption change the very market dynamics it seeks to predict?*” This maps directly onto Lopez-Lira & Tang’s empirical finding that LLM-signal alpha **decays as adoption spreads**. Surfacing this with Son/Binh shows you read both the brief *and* the literature, and it’s a genuine, publishable thesis chapter (measure signal decay against an LLM-adoption proxy) — not just a backtest. It also reframes a “will this alpha survive?” worry as a research question rather than a weakness.

7.7 Rigor guardrails — what a finance examiner will demand (where naive LLM backtests die)

- **Multiple testing / data-snooping (the #1 substance omission).** An LLM backtest has huge researcher degrees of freedom (prompt, model vintage, temperature, aggregation, horizon, universe, cost model). **Pre-register the protocol** and correct for it — Harvey, Liu & Zhu (2016, RFS); Harvey (2017, “p-hacking”); López de Prado (deflated Sharpe / backtest overfitting); White’s Reality Check / Hansen SPA. Saying this *unprompted* signals you’re past the trap that sinks most LLM-finance papers.
- **LLM look-ahead / training-cutoff contamination.** “OOS in time” $\neq$ “OOS for an LLM” — GPT-4 may *already know* how a pre-cutoff event resolved. Test strictly **post-cutoff**, show robustness across model vintages, and never feed prompt fields computed with future info. Vietnamese pre-training coverage is thin and undocumented — don’t wave it away with an English-market citation.
- **Data sufficiency (the quiet thesis-killer — do this BEFORE you commit).** Credible OOS needs enough *independent* events. VN30 $\approx$ 30 firms $\times$ 4 quarters $\times$ ~8–12 usable years, shrunk further by clean-news + post-cutoff + liquidity screens, and earnings cluster by date (cluster your SEs). **Do a back-of-envelope event count + minimum-detectable-effect calc.** If too

few post-cutoff clean events exist, the OOS-predictive framing collapses into the in-sample study you're trying to escape — know this number.

- **Transaction costs:** a flat 0.2–0.5% is a *quoted* number, not *executed*. Add a market-impact model (square-root / ADV-scaled) and report a **\$-capacity curve** — your alpha and your costs live in the same illiquid corner.
-

8. People — map of the table

- **Samuel Buerthey** — accounting lead (PEAD); the explanatory gravity, topic gatekeeper on the PEAD arm. CSR/disclosure, 0 asset-pricing.
- **Anh Tuan Nguyen** — the asset-pricing/finance method ally on *either* path. **Formally a co-sup on PEAD #247, NOT on #192** — but he co-authors the Bitcoin paper with Son + Binh, so grafting him onto #192 is natural, not a cold ask. PhD Finance, UT Arlington. Junior (h-3, ~26 cites) — can co-supervise method, can't anchor alone or open doors. **Administratively he likely can't be the registered primary supervisor either** (seniority/completions bar) — so the principal must be Binh or another senior, with Anh as associate.
- **Hà Xuân Sơn** — the LLM/AI engine; CS/security (~1.7k cites), not finance. Co-leads #192.
- **AP Binh Nguyen Thanh** — #192's *senior* name and likely its door-opener: genuine Associate Professor (h-11, ~1.1k cites; PhD in finance/economics/econometrics, Regensburg), Associate Head (Research) of The Business School, Fintech-Crypto Hub co-founder. Strong **econometrics + crypto-return-prediction + institutional weight** — but his record is **crypto / macro-uncertainty / AI-in-education, NOT equity asset-pricing**, in interdisciplinary/economics outlets (no finance-list journals). He gives #192 seniority and methods credibility, **but not the asset-pricing prong — that remains Anh's to fill**.

Net: #192's own roster is just **Binh (senior) + Son (AI)** — it has **no asset-pricing voice of its own**, and Anh is on the *PEAD* team, not #192. The ask is therefore to **graft Anh onto #192** as finance co-sup — feasible precisely because the Bitcoin paper already binds **Anh + Son + Binh** around your topic. Your job on top of that is to add the *predictive, after-cost* lens. (Binh also personifies RMIT's house-gravity toward crypto/policy over equity alpha — useful to keep in mind, never to say.)

9. Diplomacy rails (the delicate part — apply on every branch)

- ☒ **Don't pivot to Son/LLM before Buerthey answers the probe.** Naming Son's project as *shared ground* in the open is fine and planned; *redirecting* to it before the probe is the rude version — it reads as ranking the lead's project second in his own meeting. The pivot must *follow from* his answer (branch B).
- ✗ Never voice the FoR-code critique ("your project is 40% organisational behaviour") — private diagnostic, not for the team that wrote it.
- ✗ Never frame either project as "flawed." Frame as *fit between one coherent interest and two adjacent projects*.
- ✗ Never say "fallback" out loud.

- × Never ask “will Anh own the method instead of you?” — insults the lead in his own meeting (esp. VN academic hierarchy). Ask the **splitt**; the answer is the diagnostic.
 - × Don’t push back in the room if Buerthey claims all empirical framing — note it silently.
 - ☑ Address asset-pricing / method questions to **Anh** — signals you know who the finance brain is.
 - ☑ **Leave with PEAD intact**, even if you’re aiming for LLM/DR203.
-

10. On-the-spot decision rule

- **PEAD GREEN** (predictive + Anh owns method) → take PEAD; LLM = backup.
 - **PEAD EXPLANATORY AND** Son open to **LLM-under-DR203 + Anh co-sup + genuine OOS evaluation** → that’s your primary (the only configuration passing all prongs).
 - **Neither lands** (PEAD explanatory AND LLM blocked/in-sample) → no RMIT VN project passes the predictive+finance test → seriously weigh other venues for a QR-track PhD.
-

11. Likely questions to you — have answers ready

- “**Why a PhD, coming from a trading desk?**” → to formalise and publish signal research and move into a quantitative-research track; the PhD is the research training + credential. Honest and specific, not “I love academia.”
- “**What’s your NLP / LLM experience?**” → *Don’t overclaim*. None in production NLP yet — but a strong quant/ML foundation (MSc Applied Maths capstone: boosting / expert-weighting online-learning on VN equities) that ramps fast into LLM sentiment. The fit line you already sent says “quant/markets side,” so stay consistent.
- “**Why this project / why RMIT VN?**” → alumnus; direct VN-market access and data; the team’s *in-print* GenAI-sentiment→returns program (Anh + Son + Binh). (*Don’t lean on “it’s funded” — that likely means full-time only; see below.*)
- “**What would your first study be?**” → the VN30 OOS sentiment/PEAD predictive test — give the §6 scope answer (long-only/VN30 futures, net-of-cost, walk-forward, FiinPro PIT data, benchmark vs ARIMA/factor/LSTM).
- “**Full-time, or part-time alongside work?**” → *Be upfront*: you intend to **keep your full-time quant role at HSC** and do the PhD **part-time** (a friend is doing exactly this at RMIT VN — but verify her setup: part-time enrolment? self-funded or scholarship? employer support?). Three things to surface honestly:
 - **Funding tension (the catch)**: RMIT HDR stipends require *full-time* candidature and cap outside work at ~8 hrs/week — a full-time HSC job is **categorically incompatible** with holding the stipend. So keeping the job most likely means **part-time, self-funded — forgoing the stipend entirely** (partial part-time stipends are rare and need special grounds). **Mentally delete “Funded” from your target** and confirm directly.
 - **Frame the job as an asset, not a distraction**: desk-level data access, live market intuition, and a directly-relevant quant role make you an *industry-embedded* candidate — many supervisors value that *if* output holds. Say it that way.
 - **Know the cost**: part-time roughly doubles duration (~6–8 yrs) and some supervisors won’t take part-time students. Find out early — it’s a bigger dealbreaker than any framing mismatch.

(Both close 2027, but RMIT VN runs *fixed intake windows* + months of processing — ask the next intake & application dates; don't assume start timing is open.)

- “**How do you handle VN data limitations?**” → \$7.5 realities (no shorting, price limits, foreign room, survivorship/PIT) — demonstrating you've thought about feasibility is itself a strong signal.
 - “**Are you looking at other projects too?**” → Honest and disarming: “*These two adjacent projects from your team — yours and Son's — are exactly the space I care about; that's partly why I was glad Son and Anh were on the thread.*” (MAS is closed — don't volunteer it.)
 - **What admission actually requires (know before you commit):** you must **secure a supervisor first** (admission is conditional on it); submit **your own research proposal** (the project text is *their* pitch, not your proposal — your \$6 scope is raw material); pass **Confirmation of Candidature** (~12 mo full-time / ~24 mo part-time — a written + oral gate that can end candidature); and confirm your **English-waiver recency** (an RMIT-VN MSc taught in English usually qualifies, but waivers can expire ~2 yrs). Tuition: confirm the **part-time local fee** and whether any **fee waiver can attach to part-time** even when the stipend can't.
 - **Conflict of interest / data / IP (raise it before they do):** you trade VN equities at HSC and propose a tradable VN-equity thesis. Pre-empt the obvious worry — clarify HSC **data licensing**, **no use of MNPI**, and whether your **employer has any IP claim** on the research. Naming it yourself reads as professionalism; being caught out by it reads as naïveté.
-

12. Stepping back — the questions only you can decide (reflect AFTER the meeting)

The brief optimises *which project, framed how*. These matter more for a QR career, and no probe can answer them for you:

- **Can these specific people train me to publish where quants care?** A QR career is built on publications + methods training + a name that opens doors — far more than the project label. Honest read of the table: **Binh** = real econometrics + seniority, but not asset-pricing (and not in the room); **Anh** = the only true finance voice, but junior (h-3), can't open doors alone; **Son** = CS. A workable *methods* team **if** they allow predictive work — but no top-tier asset-pricing name who lands you at a fund on reputation. Weigh that soberly.
 - **Do I actually want to work with them for 6–8 years?** (Part-time horizon.) Supervision fit — responsiveness, feedback style, trust — decides whether a PhD finishes or quietly dies. Judge the humans, not just the topic.
 - **Is RMIT VN the right venue at all?** A teaching-heavy offshore campus — not disqualifying, but go in clear-eyed that you may be trading the strongest research training for funding/convenience. Branch C (“weigh other venues”) deserves more weight than this document's momentum implies.
 - **Does part-time-while-working actually work here?** (See §11.) Settle the funding + workload + supervisor-acceptance reality before committing — it's the practical dealbreaker.
-

Appendix A — full project descriptions (verbatim, from the RMIT VN catalogue)

Source: data/rmit_vn_2026-06-17.json. Reference text so you can speak to the exact framing; the distilled version is in §1.5.

A.1 — LLM #192: *Large Language Models for Business Innovation and Decision-Making*

Financial markets aggregate vast quantities of information — corporate disclosures, macro-economic data, news, and investor sentiment — into asset prices. Understanding how efficiently this process occurs, and how it can be supported or improved through AI, is one of the central questions of modern financial economics. This project investigates whether and how Large Language Models (LLMs) can advance our ability to predict asset prices, characterise investor behaviour, and assess the informational efficiency of financial markets.

Traditional asset price prediction models — from ARIMA and factor models to early deep learning approaches — have demonstrated predictive power in narrow conditions but struggle with the non-stationarity, regime changes, and behavioural dynamics that characterise real markets. LLMs offer a qualitatively different capability: the ability to process and interpret the unstructured textual information — analyst commentary, earnings call transcripts, news sentiment — that drives investor expectations and, ultimately, asset prices.

The efficient market hypothesis (EMH) holds that asset prices reflect all available information. Yet decades of empirical research in behavioural finance have documented systematic deviations from this ideal — momentum effects, overreaction and underreaction, herding, and sentiment-driven mispricing — that suggest markets are not fully efficient and that investor psychology plays a meaningful role in price formation.

This tension between market efficiency and behavioural anomalies is precisely where LLMs may offer the most valuable contributions. If asset prices are partly driven by how investors interpret and respond to information — rather than by the information itself — then AI systems that can model language, sentiment, and narrative at scale may capture predictive signals that traditional quantitative models miss. At the same time, the widespread adoption of LLM-based trading systems raises a deeper question: does AI adoption change the very market dynamics it seeks to predict? If LLMs become pervasive tools of investment analysis, their influence on investor behaviour and price formation could alter the efficiency properties of markets in ways that are not yet understood.

This project is motivated by the conviction that these are not merely technical questions but fundamental financial economics questions about the nature of information, the behaviour of investors, and the efficiency of markets in an AI-augmented era. The findings will be relevant to investors, regulators, and policymakers seeking to understand the implications of AI adoption for market structure and financial stability.

Team: Hà Xuân Sơn, AP. Binh Nguyen · *Discipline:* Business (GSBL) · *Codes:* DR201 / DR205 / DR203 · *FoR:* 150311 Organisational Behaviour 40% / 080109 Pattern Recognition & Data Mining 35% / 150310 Innovation & Technology Management 25% · *Funded, closes 30 Jun 2027.*

A.2 — PEAD #247: *Investor Sentiment, Earnings Surprises, and Market Anomalies Leveraging Generative AI in Capital Market Research*

This project aims to integrate behavioral finance with AI-driven sentiment analysis to investigate how investor sentiment influences market reactions to earnings surprises, and whether it amplifies or mitigates anomalies such as post-earnings announcement drift (PEAD).

The relationship between accounting information and capital markets remains a central theme in financial accounting research. A key focus has been how investors interpret and respond to earnings announcements, particularly in the context of PEAD. Recent work highlights investor attention and sentiment as important mechanisms influencing how markets process such information. With the rise of digital platforms, sentiment now spreads rapidly through social media and online news, increasingly shaping trading behavior.

At the same time, advances in large language models (LLMs), such as ChatGPT, offer powerful new tools for extracting sentiment from unstructured online data, complementing traditional proxies like Google Trends or media coverage. This project will employ LLMs to analyze sentiment from digital sources and assess whether these AI-powered measures explain investor overreaction or underreaction to earnings surprises. It will also test whether such sentiment dynamics intensify or dampen anomalies like PEAD.

By bridging behavioral finance with frontier methods in generative AI, this study seeks to advance understanding of how sentiment, attention, and accounting information jointly influence price discovery and asset pricing in modern financial markets.

Team: Samuel Buerthey, Anh Tuan Nguyen, Ha Xuan Son · *Discipline:* Accounting (AISSC) · *Code:* DR200 · *FoR:* 350202 Financial Accounting 40% / 350103 Finance 30% / 460299 Artificial Intelligence 30% · *Funded, closes 31 Dec 2027.*